

Disco/Hadoop MapReduce

Benjamin Zaitlen
Continuum Analytics

PyCon 2013



What is Large Data?

What is Large Data?

- Can't fit into Excel

What is Large Data?

- Can't fit into Excel
 - Increase Memory

What is Large Data?

- Can't fit into Excel
 - Increase Memory
- Can't fit into R

What is Large Data?

- Can't fit into Excel
 - Increase Memory
- Can't fit into R
 - Increase Memory

What is Large Data?

- Can't fit into Excel
 - Increase Memory
- Can't fit into R
 - Increase Memory
- Can't fit into Memory

What is Large Data?

- Can't fit into Excel
 - Increase Memory
- Can't fit into R
 - Increase Memory
- Can't fit into Memory
 - Increase Memory

What is Large Data?

- Can't fit into Excel
 - Increase Memory
- Can't fit into R
 - Increase Memory
- Can't fit into Memory
 - Increase Memory
- Can't fit on a single disk

What is Large Data?

- Can't fit into Excel
 - Increase Memory
- Can't fit into R
 - Increase Memory
- Can't fit into Memory
 - Increase Memory
- Can't fit on a single disk
 - Distributed Filesystem: SAN, HDFS/DDFS, AWS: S3, Redshift, etc.

MapReduce

Framework to help solve the problem of distributed computation for distributed data



MapReduce

Framework to help solve the problem of distributed computation for distributed data

- A mass of data: records

MapReduce

Framework to help solve the problem of distributed computation for distributed data

- A mass of data: records
- Split/**Map** records into key-values pairs

MapReduce

Framework to help solve the problem of distributed computation for distributed data

- A mass of data: records
- Split/**Map** records into key-values pairs
- Collect/Partition kv pairs (Optional Sort)

MapReduce

Framework to help solve the problem of distributed computation for distributed data


- A mass of data: records
- Split/**Map** records into key-values pairs
- Collect/Partition kv pairs (Optional Sort)
- Buckets are passed to **Reduce** function

MapReduce


Framework to help solve the problem of distributed computation for distributed data

- A mass of data: records
- Split/**Map** records into key-values pairs
- Collect/Partition kv pairs (Optional Sort)
- Buckets are passed to **Reduce** function
- Result is returned

MapReduce Workflow




MapReduce Workflow



- Push code to data

MapReduce Workflow



- Push code to data
- Lots of network traffic

MapReduce it's a Party



MapReduce it's a Party



MapReduce it's a Party



MapReduce it's a Party



- Disco: Python + Erlang

MapReduce it's a Party



- Disco: Python + Erlang
- Hadoop: Java-Dumbo (Python Streaming)

MapReduce it's a Party



- Disco: Python + Erlang
- Hadoop: Java-Dumbo (Python Streaming)
- Distributed FileSystem: DDFS

MapReduce it's a Party



- Disco: Python + Erlang
- Hadoop: Java-Dumbo (Python Streaming)
- Distributed FileSystem: DDFS
- Bring your friends
 - NumPy
 - SciPy
 - pandas
 - scikits-learn
 - OpenCV

Canonical Example

```
1
2 from disco.job import Job
3 from disco.worker.classic.func import chain_reader
4 from disco.core import result_iterator
5
6 class WordCount(Job):
7     partitions = 2
8     input=["sherlock_complete.txt","poirot_complete.txt"] #collected works
9
10    @staticmethod
11    def map(line, params):
12        import string
13        for word in line.split():
14            yield strippedWord, 1
15
16    @staticmethod
17    def reduce(iter, params):
18        from disco.util import kvgroup
19        for word, counts in kvgroup(sorted(iter)):
20            yield word, sum(counts)
21
22 if __name__ == "__main__":
23     from MapReduce_CountWords_Chain import WordCount
24
25     for (word, counts) in result_iterator(WordCount.wait(show=False)):
26         print word, counts
```


Canonical Example

```
1
2 from disco.job import Job
3 from disco.worker.classic.func import chain_reader
4 from disco.core import result_iterator
5
6 class WordCount(Job):
7     partitions = 2
8     input=["sherlock_complete.txt","poirot_complete.txt"] #collected works
9
10    @staticmethod
11    def map(line, params):
12        import string
13        for word in line.split():
14            yield strippedWord, 1
15
16    @staticmethod
17    def reduce(iter, params):
18        from disco.util import kvgroup
19        for word, counts in kvgroup(sorted(iter)):
20            yield word, sum(counts)
21
22 if __name__ == "__main__":
23     from MapReduce_CountWords_Chain import WordCount
24
25     for (word, counts) in result_iterator(WordCount.wait(show=False)):
26         print word, counts
```


Canonical Example

```
1
2 from disco.job import Job
3 from disco.worker.classic.func import chain_reader
4 from disco.core import result_iterator
5
6 class WordCount(Job):
7     partitions = 2
8     input=["sherlock_complete.txt","poiroit_complete.txt"] #collected works
9
10    @staticmethod
11    def map(line, params):
12        import string
13        for word in line.split():
14            yield strippedWord, 1
15
16    @staticmethod
17    def reduce(iter, params):
18        from disco.util import kvgroup
19        for word, counts in kvgroup(sorted(iter)):
20            yield word, sum(counts)
21
22 if __name__ == "__main__":
23     from MapReduce_CountWords_Chain import WordCount
24
25     for (word, counts) in result_iterator(WordCount.wait(show=False)):
26         print word, counts
```


Chain Jobs




Chain Jobs



Chain Jobs



Chain Jobs



MapReduce Thoughts



MapReduce Thoughts

- Data Cleansing
 - Everyone's pain point

MapReduce Thoughts


- Data Cleansing
 - Everyone's pain point
- Task Deconstruction
 - Good for code management
 - Hides -- in a good way -- data management


MapReduce Thoughts

- Data Cleansing
 - Everyone's pain point
- Task Deconstruction
 - Good for code management
 - Hides -- in a good way -- data management
- Can Be Inefficient
 - Network traffic
 - Job organization

Data Thoughts





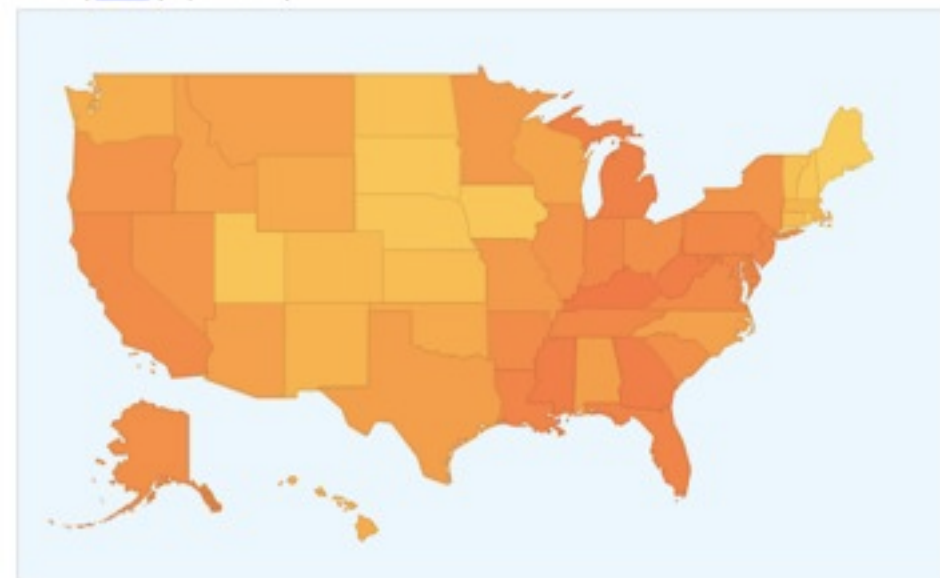



83 percent of the
radiologists
missed the gorilla

Google Flu


Explore flu trends - United States

We've found that certain search terms are good indicators of flu activity. Google Flu Trends uses aggregated Google search data to estimate flu activity. [Learn more »](#)



- Data Mining
- Faster than CDC

Google Get's Wrong



- Typically, prediction is great!
- This year not so much
- Google: No comment!
- Feedback mechanism from hype-up media

Wrapping Up

- Invisible Gorillas will stay invisible
- Turnkey analytics is dangerous
- Good Analysis
 - Requires iterative exploration
 - Peer review and collaboration